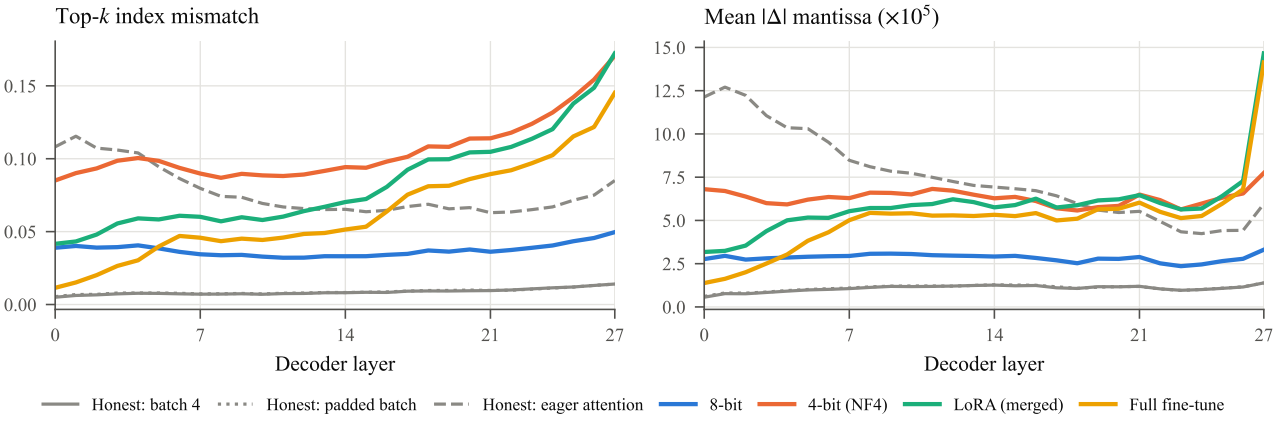
Top-k index mismatch
0.15
0.10
0.05
0.00
0
7
14
21
27
Decoder layer
Mean |Δ| mantissa (×10^5)
15.0
12.5
10.0
7.5
5.0
2.5
0.0
0
7
14
21
27
Decoder layer
Honest: batch 4
Honest: padded batch
Honest: eager attention
8-bit
4-bit (NF4)
LoRA (merged)
Full fine-tune